

NO.1 Check the image version in pod without the describe command

Answer:

```
kubectl get po nginx -o  
jsonpath='{.spec.containers[].image}'{"\n"}
```

NO.2 Create a nginx pod with label env=test in engineering namespace

Answer:

```
kubectl run nginx --image=nginx --restart=Never --labels=env=test --namespace=engineering  
--dry-run -o yaml > nginx-pod.yaml kubectl run nginx --image=nginx --restart=Never --labels=env=test  
--namespace=engineering --dry-run -o yaml | kubectl create -n engineering -f - YAML File:  
apiVersion: v1  
kind: Pod  
metadata:  
name: nginx  
namespace: engineering  
labels:  
env: test  
spec:  
containers:  
- name: nginx  
image: nginx  
imagePullPolicy: IfNotPresent  
restartPolicy: Never  
kubectl create -f nginx-pod.yaml
```

NO.3 Score: 4%



Task

Create a persistent volume with name app-data , of capacity 1Gi and access mode ReadOnlyMany. The type of volume is hostPath and its location is /srv/app-data .

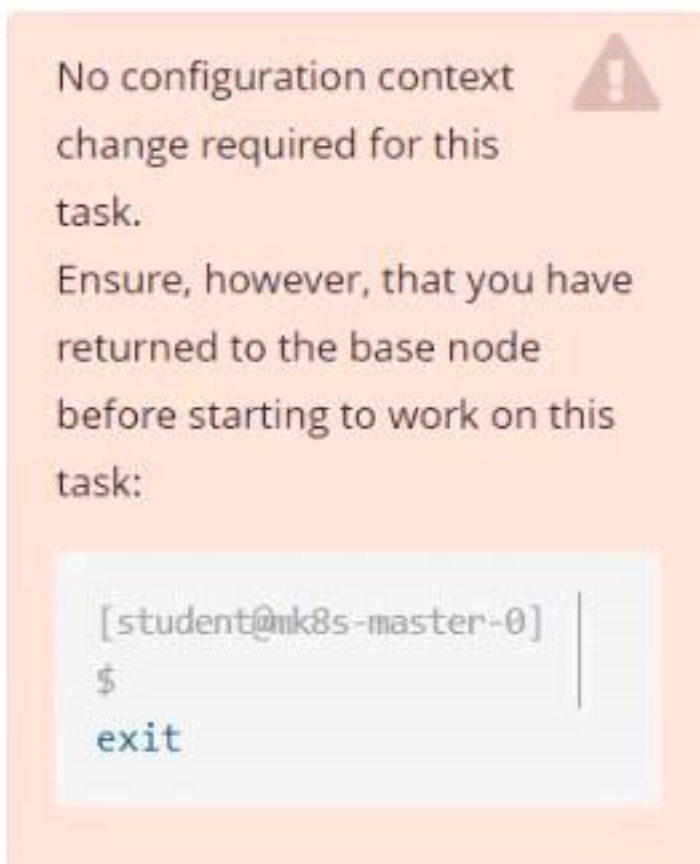
Answer:

Solution:

```
#vi pv.yaml  
apiVersion: v1  
kind: PersistentVolume
```

```
metadata:
name: app-config
spec:
capacity:
storage: 1Gi
accessModes:
- ReadOnlyMany
hostPath:
path: /srv/app-config
#
kubectl create -f pv.yaml
```

NO.4 Score: 7%



Task

First, create a snapshot of the existing etcd instance running at <https://127.0.0.1:2379>, saving the snapshot to /srv/data/etcd-snapshot.db.

Creating a snapshot of the given instance is expected to complete in seconds. If the operation seems to hang, something's likely wrong with your command. Use **CTRL + C** to cancel the operation and try again.

Next, restore an existing, previous snapshot located at `/var/lib/backup/etcd-snapshot-previous.db`

The following TLS certificates/key are supplied for connecting to the server with `etcdctl`:

- CA certificate:
`/opt/KUIN00601/ca.crt`
- Client certificate:
`/opt/KUIN00601/etcd-client.crt`
- Client key:
`/opt/KUIN00601/etcd-client.key`

Answer:

Solution:

#backup

```
ETCDCTL_API=3 etcdctl --endpoints="https://127.0.0.1:2379" --cacert=/opt/KUIN000601/ca.crt  
--cert=/opt/KUIN000601/etcd-client.crt --key=/opt/KUIN000601/etcd-client.key snapshot save  
/etc/data/etcd-snapshot.db
```

#restore

```
ETCDCTL_API=3 etcdctl --endpoints="https://127.0.0.1:2379" --cacert=/opt/KUIN000601/ca.crt  
--cert=/opt/KUIN000601/etcd-client.crt --key=/opt/KUIN000601/etcd-client.key snapshot restore  
/var/lib/backup/etcd-snapshot-previoys.db
```

NO.5 Score: 7%



Task

Create a new nginx Ingress resource as follows:

- * Name: ping
- * Namespace: ing-internal
- * Exposing service hi on path /hi using service port 5678



Answer:

Solution:

```
vi ingress.yaml
```

```
#
```

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
name: ping
namespace: ing-internal
spec:
rules:
- http:
paths:
- path: /hi
pathType: Prefix
backend:
service:
name: hi
port:
number: 5678
#
kubectl create -f ingress.yaml
```

NO.6 Create a persistent volume with name app-data, of capacity 2Gi and access mode ReadWriteMany. The type of volume is hostPath and its location is /srv/app-data.

Answer:

solution

Persistent Volume

A persistent volume is a piece of storage in a Kubernetes cluster. PersistentVolumes are a cluster-level resource like nodes, which don't belong to any namespace. It is provisioned by the administrator and has a particular file size. This way, a developer deploying their app on Kubernetes need not know the underlying infrastructure. When the developer needs a certain amount of persistent storage for their application, the system administrator configures the cluster so that they consume the PersistentVolume provisioned in an easy way.

Creating Persistent Volume

```
kind: PersistentVolume apiVersion: v1 metadata: name:app-data spec: capacity: # defines the
capacity of PV we are creating storage: 2Gi #the amount of storage we are trying to claim
accessModes: # defines the rights of the volume we are creating - ReadWriteMany hostPath: path:
"/srv/app-data" # path to which we are creating the volume Challenge Create a Persistent Volume
named app-data, with access mode ReadWriteMany, storage classname shared, 2Gi of storage
capacity and the host path /srv/app-data.
```

```

apiVersion: v1
kind: PersistentVolume
metadata:
  name: app-data
spec:
  capacity:
    storage: 2Gi
  accessModes:
    - ReadWriteMany
  hostPath:
    path: /srv/app-data
  storageClassName: shared
~
~
~
~
~
~
~
~
~
~
"app-data.yaml" 12L, 194C

```

2. Save the file and create the persistent volume.

```

njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl create -f pv.yaml
persistentvolume/pv created

```

3. View the persistent volume.

```

njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl get pv

```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
app-data	2Gi	RWX	Retain	Available		shared		31s

Our persistent volume status is available meaning it is available and it has not been mounted yet. This status will change when we mount the persistentVolume to a persistentVolumeClaim.

PersistentVolumeClaim

In a real ecosystem, a system admin will create the PersistentVolume then a developer will create a PersistentVolumeClaim which will be referenced in a pod. A PersistentVolumeClaim is created by specifying the minimum size and the access mode they require from the persistentVolume.

Challenge

Create a Persistent Volume Claim that requests the Persistent Volume we had created above. The claim should request 2Gi. Ensure that the Persistent Volume Claim has the same storageClassName as the persistentVolume you had previously created.

```
kind: PersistentVolume apiVersion: v1 metadata: name:app-data
```

```
spec:
```

```
accessModes: - ReadWriteMany resources:
```

```
requests: storage: 2Gi
```

```
storageClassName: shared
```

2. Save and create the pvc

```
njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl create -f app-data.yaml
```

```
persistentvolumeclaim/app-data created
```

3. View the pvc

```

njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl get pvc

```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS
pv	Bound	pv	512m	RWX	shared

4. Let's see what has changed in the pv we had initially created.

```
njerry191@cloudshell:~ (extreme-clone-265411)$ kubectl get pv
NAME      CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM          STORAGECLASS  REASON  AGE
pv        512m      RWX           Retain          Bound   default/pv     shared       16m
```

Our status has now changed from available to bound.

5. Create a new pod named myapp with image nginx that will be used to Mount the Persistent Volume Claim with the path /var/app/config.

Mounting a Claim

```
apiVersion: v1 kind: Pod metadata: creationTimestamp: null name: app-data spec: volumes: -
name: configpvc persistenVolumeClaim: claimName: app-data containers: - image: nginx name: app
volumeMounts: - mountPath: "/srv/app-data " name: configpvc
```

NO.7 List "nginx-dev" and "nginx-prod" pod and delete those pods

Answer:

```
kubect1 get pods -o wide
```

```
kubectl delete po "nginx-dev" kubectl delete po "nginx-prod"
```